

High Performance Python Lab

Overview of multithreading and multiprocessing tools

November 22, 2023

Difference between multithreading and multiprocessing

Multithreading:

- + **less overhead**
- **common address space**
- **need to synchronize access to resources**
- + **easy to communicate with others**

Multiprocessing:

- **more overhead**
- **independent address spaces for each process**
- + **less need to synchronize**
- **need to find ways to communicate with others**

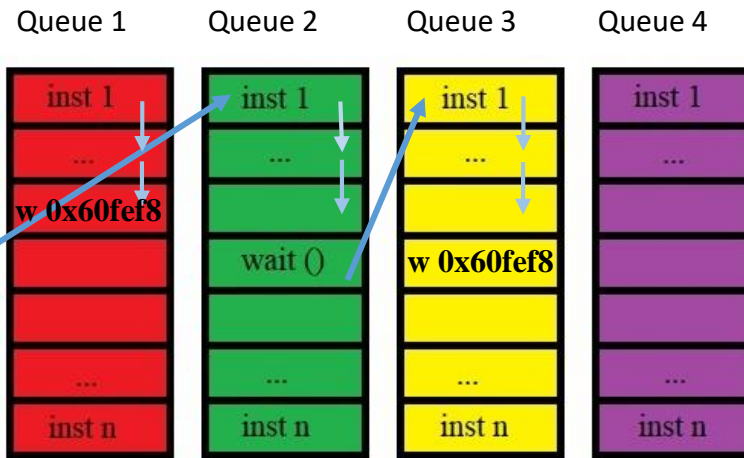
Overview of libraries

- Multithreading
- Multiprocessing
- Concurrent
- Joblib



Also: Dask, ray, asyncio

Synchronization


Threads from software point of view



Threads from hardware point of view



Synchronization



Synchronization techniques

Busy waiting – repeatedly check some condition. Once it's true you can modify some shared variable.

Mutex – a technique of allowing only one thread to access a variable at a given point in time using locks. Once a thread **acquires the lock** the shared variable cannot be accessed by other threads.


Semaphore – a generalization of mutex. When acquiring the lock semaphore decrements the counter. Once the counter gets to zero other threads have to wait while others unlock the semaphore.

Model examples

Pipeline – if process can be broken down into stages of a single pipeline independent parts can be done in parallel. Or they can be done in parallel on different data (**query processing, web server administration**).

Background Task – a thread running in the background doing "menial" jobs (**back up data, collect garbage**).

User interface – GUI use multithreading to ensure responsiveness of the interface to the user (**progress bars, loading animation**).



Semaphore use case description

Web application:

- it has to process incoming **requests** from users. Seeing as these requests are independent the **processing** can be done simultaneously i.e using **threads**.
- there are **millions of requests** and the server can only provide a **hundred of threads**. There is no way this number of threads can deal with such number of requests in a runtime fashion.
- that is why the requests need to be **queued up**. To do that the servers have **queuing threads**.
- the length of the queue is **limited by the memory**.


Implementation

The model for the case is called **Consumer-Producer model**. And this is a typical use case for semaphores.

What do we need?

Implementation

The model for the case is called **Consumer-Producer model**. And this is a typical use case for semaphores.




What do we need?

1) Queue – to serialize request processing

Implementation

The model for the case is called **Consumer-Producer model**. And this is a typical use case for semaphores.




What do we need?

- 1) Queue – to serialize request processing
- 2) Producer/s – thread/s that try to fill the queue

Implementation

The model for the case is called **Consumer-Producer model**. And this is a typical use case for semaphores.




What do we need?

- 1) Queue – to serialize request processing
- 2) Producer/s – thread/s that try to fill the queue
- 3) Consumer/s - thread/s that try to empty the queue

Implementation

The model for the case is called **Consumer-Producer model**. And this is a typical use case for semaphores.




What do we need?

- 1) Queue – to serialize request processing
- 2) Producer/s – thread/s that try to fill the queue
- 3) Consumer/s - thread/s that try to empty the queue

??? Synchronization ???

Implementation

The model for the case is called **Consumer-Producer model**. And this is a typical use case for semaphores.



What do we need?

- 1) Queue – to serialize request processing
- 2) Producer/s – thread/s that try to fill the queue
- 3) Consumer/s - thread/s that try to empty the queue

??? Synchronization ???

- 1) Semaphore to protect queue size going over capacity
- 2) Semaphore to protect queue size going negative

Global interpreter lock (GIL)

Python code is interpreted line by line. In a multithreaded program the interpreter is **locked** by a single thread. This is needed to avoid segmentation faults.

Hence multithreading can't properly help with performance if the code needs constant interpretation – **CPU bound programs**.

Alternatively, if the program needs significant amount of time to carry out the already interpreted instruction (**I/O calls**, for instance) then GIL less of an issue

Practice session

- 1) Introductory examples
- 2) Multithreading/multiprocessing with Python on more realistic use cases
- 3) Little bit of practice

Practice session

- 1) Introductory examples
- 2) Multithreading/multiprocessing with Python on more realistic use cases
- 3) Little bit of practice

Conclusions:

- 1) If the program/function occasionally lets go of the CPU then GIL is less of a problem
- 2) Multiprocessing adds time overhead on creating and destroying processes
- 3) Multithreading and multiprocessing add memory usage load
- 4) Popular libraries (scikit, torch, numpy etc.) usually have multiprocessing/multithreading modules already implemented

Multiprocessing/multithreading inside libraries

Numpy: - Blas and Lapack via Cython

- ...



Pytorch: - OpenMP and TBB via Cython

- JIT compilation

- ...



Scikit: - joblib

- ...