


Final Remarks on Planning: Learning and Beyond

Planning Algorithms in AI

How to choose a Planning algorithm?

First an most important question is to think about the **state space** and **action space**.

Can the state space be discretized?




Yes, and the dimensionality remains relatively low (3-4)

Discrete planning (L02)

- Forwards search
- Dijkstra, A* if cost

Yes, but dimensions are high

No, irregular regions, complicated spaces, etc.



Sampling (L04) (LaValle)

- Variants of RRT or PRM
- RRT* if we want optimality.

How to choose a Planning algorithm?



Do we want a single plan or a **multiple query** algorithm?

Yes, plan can be as complex as we want



**Discrete planning (L02),
RRT (L04),
Traj. Optimization (L06)**

No, we will execute the plan multiple times

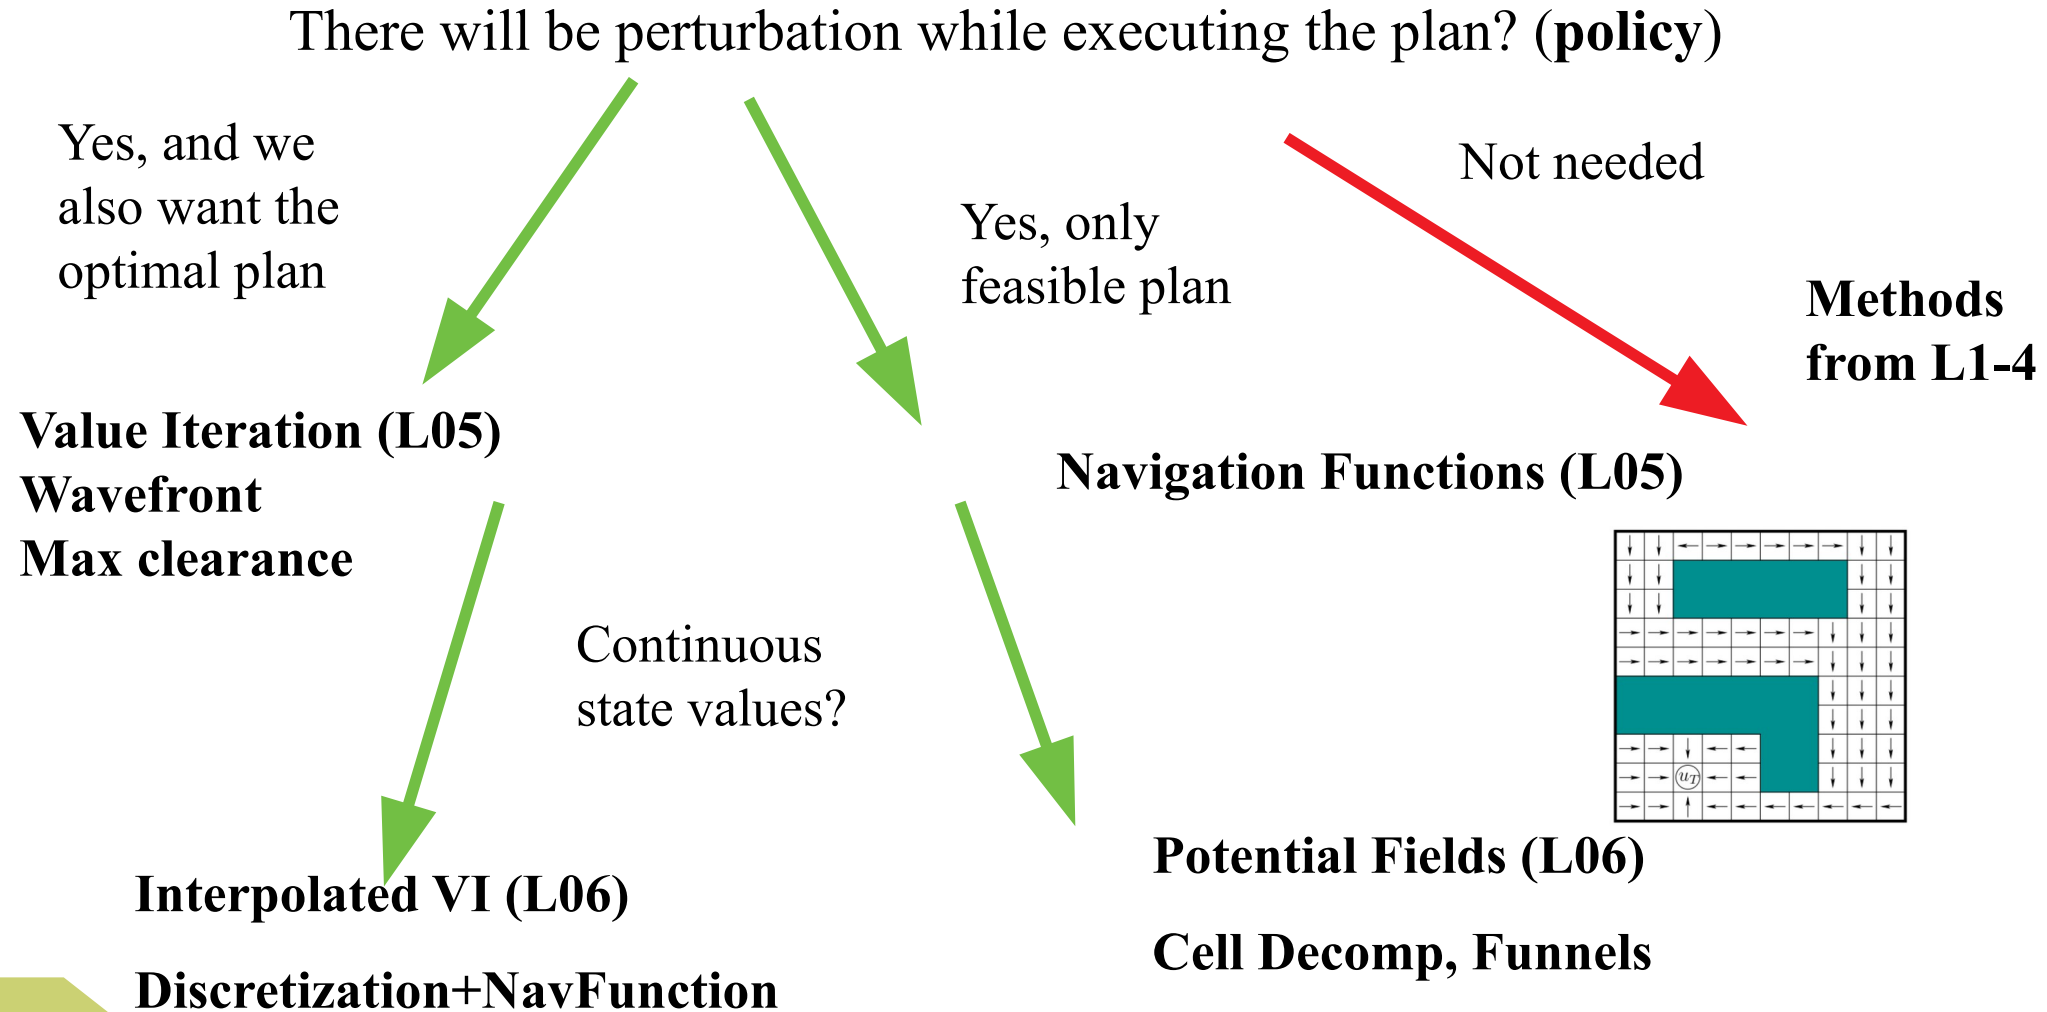


PRM (L04)

Discrete Planning (L02) (it might be still more efficient)

Multiple-VI (L05) (only if memory is not a problem)


How to choose a Planning algorithm?



How to choose a Planning algorithm?

Is this a game? → gigantic state space


State space is
“small” enough for
building a decision
tree



Alpha-Beta Pruning (Russel&Norvig)


It allows to almost double the depth wrt naive decision tree and the solutions are guaranteed to be optimal

State space is monstrous
but the action space is
relatively small



Monte-Carlo Tree Search

The evaluation is approximated by sampling and this allows to explore depth of the tree unthinkable for exhaustive search



How to choose a Planning algorithm?

Is there **uncertainty** of any kind?

Yes

MDP (L08) (Sutton&Barto)

- It works well for finite state spaces
- Requires known transitions

Yes, but the environment might be unknown

Reinforcement Learning

- It only interacts with the environment and does not require **complete** knowledge.


No, states are observable and transitions are deterministic

Planning methods from L1-6

How to choose a Planning algorithm?

Are any parameters of the environment **unknown**?

Yes, such as
transition function,
action execution,
etc.



Learning Based approaches:

- **Supervised learning**
- **Reinforcement Learning**

No, we have total control
on all elements



**Planning L1-4, discrete,
sampling, policies, etc.**


MDP (L08) For discrete and low
dimensional state spaces




Dimensionality is high, we need
function NN approximators

Planning Taxonomy

Now we see the power of combining plans and a hierarchical approach.




Open loop execution



Feedback



Refinement



Hierarchical

Learning in planning: Limits and Potentials

- Some problems do not require learning and planning methods can be enough, so it is better if we ask first: does it make sense for your problem?
- We must always think what is the learning algorithm doing:
 - Policy approximated by a parametrized function
 - Compressing a value function over a continuous state space.
 - Heuristics for a distance function
- Second we must think about the state space:
 - Position, Pose, Image, a sequence of values, etc.
- The equivalent in CV to classification and regression is not always applicable since the agent, as a result of a plan, executes an action in an environment, whose results are uncertain as well.

Learning in planning: Limits and Potentials

- On the other hand, learning-based approaches **allow** to solve planning under impossible conditions for planning, such as unknown dynamics, changes in parameters and many other artifacts.
- **Combining** both approaches can be a great way to go. Examples are MCTS with learned policies, which also combines approximation to evaluation functions.
 - In RRT, learning sampling distributions, A* learning heuristics, Potential functions learned, etc.