

Models with long memory

Skoltech

Alexey Zaytsev

Some slides are from Zabolotnyi Artem

Benchmarking for Efficient Transformers for transactions data

Dataset description

Transaction dataset

Task

Predict the probability of user default

Features

~10 categorical features (MCC code, type etc)

~10 numerical (amount, time etc)

Train test split

Train – 1 million users

Validation – 250k users

Out-of-time – 500k users


Length

700 - mean transaction count

(*Typical transformer length 512*)

Metric

Gini coefficient



Hyperparameters search for Longformer

Quality

Longformer model shows better quality than baseline LSTM


Hyperparameters

- Very deep model required a huge amount of memory for training
- Sliding window size W does not change quality significant


Baseline metric

LSTM - 60.12 ± 0.17


| № | Heads | Embedding size | W | Layers | Hidden dropout | OOT Gini | RAM (Mb) |
|----|-------|----------------|-----|--------|----------------|--------------|-------------|
| 1 | 8 | 512 | 128 | 1 | 0.1 | 59.59 | 7647 |
| 2 | 8 | 256 | 128 | 1 | 0.1 | 58.84 | 4929 |
| 3 | 8 | 64 | 128 | 1 | 0.1 | 60.11 | 4581 |
| 4 | 8 | 64 | 128 | 2 | 0.1 | 61.02 | 6039 |
| 5 | 8 | 64 | 256 | 2 | 0.1 | 61.03 | 10131 |
| 6 | 8 | 64 | 64 | 2 | 0.1 | 60.58 | 4500 |
| 7 | 4 | 64 | 128 | 8 | 0.3 | 62.21 | 13971 |
| 8 | 4 | 32 | 128 | 8 | 0.3 | 62.58 | 13505 |
| 9 | 4 | 64 | 128 | 8 | 0.5 | 62.39 | 13971 |
| 10 | 4 | 64 | 128 | 12 | 0.3 | 62.39 | 20255 |
| 11 | 2 | 64 | 128 | 8 | 0.5 | 62.47 | 12891 |
| 12 | 4 | 64 | 128 | 12 | 0.6 | 62.83 | 20255 |




(a) Full n^2 attention



(b) Sliding window attention



(c) Dilated sliding window



(d) Global+sliding window

Results of the experiment

Matrix projection Linformer

Quality


Linformer with projection matrix shows poor results in comparing with recurrent neural network models

Shift invariance of a projection matrix

The projection matrix not being shifting invariant

Shift invariance test

We train the model to prove this property, choose a subsample, and calculate the metric on it. Then we shift all transactions by one and calculate the metric again.



Results of the experiment

Matrix projection Linformer

Quality


Linformer with projection matrix shows poor results in comparing with recurrent neural network models

Shift invariance of a projection matrix

The projection matrix not being shifting invariant

Shift invariance test

We train the model to prove this property, choose a subsample, and calculate the metric on it. Then we shift all transactions by one and calculate the metric again.



Linformer projection matrix

| Dataset | Gini |
|--------------------------|-------|
| Original dataset | 53.47 |
| Shift by one transaction | 47.31 |

Results of the experiment

Convolution Linformer

Quality


Convolution Linformer shows a better result using less memory and shorter training time.

Convolution modification

To make projection shift-invariant, we replace a projection matrix with a convolution layer.

Cycle-shift augmentation

The result of convolution with kernel and stride equals size could change after adding a new object in sequence. (+1 Gini)



Results of the experiment

Convolution Linformer

Metric

LSTM - 60.12±0.17

Longformer - 62.80

Quality

Convolution Linformer shows a better result using less memory and shorter training time.

Convolution modification

To make projection shift-invariant, we replace a projection matrix with a convolution layer.

Cycle-shift augmentation

The result of convolution with kernel and stride equals size could change after adding a new object in sequence. (+1 Gini)

| № | Projection size | Heads | Embedding size | Layers | Hidden dropout | dim ff | OOT gini | RAM (Mb) |
|----|--------------------|-------|-------------------|--------|-------------------|--------|--------------------|-------------|
| 1 | 128 | 8 | 64 | 4 | 0.2 | 2048 | 63.5 | 8757 |
| 2 | 32 | 8 | 64 | 4 | 0.2 | 2048 | 62.02 | 5124 |
| 3 | 64 | 8 | 64 | 4 | 0.2 | 2048 | 62.40 | 6789 |
| 4 | 128 | 4 | 64 | 4 | 0.2 | 2048 | 62.73 | 7159 |
| 5 | 128 | 8 | 64 | 2 | 0.2 | 2048 | 64.04 | 5333 |
| 6 | 128 | 8 | 64 | 4 | 0.2 | 2048 | 62.49 | 8715 |
| 7 | 128 | 8 | 32 | 4 | 0.2 | 2048 | 64.17 | 8179 |
| 8 | 128 | 8 | 64 | 4 | 0.2 | 2048 | 62.98 | 3635 |
| 9 | 128 | 8 | 64 | 3 | 0.2 | 2048 | 63.53 | 7051 |
| 10 | 128 | 8 | 64 | 2 | 0.2 | 512 | 64.3 ± 0.19 | 5333 |
| 11 | 128 | 8 | 64 | 2 | 0.2 | 256 | 63.5 | 3395 |
| 12 | 128 | 8 | 128 | 2 | 0.2 | 512 | 63.26 | 3833 |

Results of the experiment

Convolution Linformer

Curriculum learning: Pre-trained weights

Training first on short sequences and the fine-tune on long show better quality and reduce training time

Non linear complexity

We need to increase the projection size:
increasing length of sequence without increasing projection size leads to quality degradation.

| № | Length | K | Training type | OOT gini | RAM (Mb) | Train time (hours) |
|---|--------|-----|-----------------------------------|---------------------|-------------|-----------------------|
| 1 | 350 | 128 | From scratch | 64.37 | 9787 | 3 |
| 2 | 1500 | 128 | From scratch | 63.75 | 12451 | 4.5 |
| 3 | 1500 | 256 | From scratch | 64.81 | 16047 | 11.61 |
| 4 | 1500 | 256 | Pre-trained weights | 65.2 | 16081 | 9.15 |
| 5 | 1500 | 400 | From scratch | 65.39 | 23481 | 17.3 |
| 6 | 1500 | 400 | Pre-trained/train only conv layer | 64.05 | 2299 | 7.1 |
| 7 | 1500 | 400 | Pre-trained weights | 66.00 ± 0.21 | 23489 | 8.5 |

Results of the experiment

Comparison LSTM and Transformer

| Transaction count | LSTM | Longformer | Linformer | Convolution Linformer |
|----------------------|------------------|------------|-----------|--------------------------|
| 350 | 60.12 ± 0.17 | 62.80 | 53.47 | 64.37 ± 0.19 |
| 1500 | 63.01 ± 0.14 | None | None | 66.00 ± 0.21 |

Performer: making transformers faster


Tay, Yi, et al. "Long range arena: A benchmark for efficient transformers." *arXiv preprint arXiv:2011.04006* (2020).

Overall results

Speed (x axis), Performance (y axis) and Memory (size of the circles) for different models

BigBird is better than the vanilla transformer

Performer is more efficient



Performer

$$\exp(q_i^T k_j) = \phi(q_i)^T \phi(k_j) = \mathbb{E}(\phi(q_i)^T \phi(k_j))$$


Vanilla attention formula:

$$\text{Att}_{\leftrightarrow}(\mathbf{Q}, \mathbf{K}, \mathbf{V}) = \mathbf{D}^{-1} \mathbf{A} \mathbf{V}, \quad \mathbf{A} = \exp(\mathbf{Q} \mathbf{K}^\top / \sqrt{d}), \quad \mathbf{D} = \text{diag}(\mathbf{A} \mathbf{1}_L).$$

Efficient bidirectional attention:

$$\widehat{\text{Att}}_{\leftrightarrow}(\mathbf{Q}, \mathbf{K}, \mathbf{V}) = \widehat{\mathbf{D}}^{-1} (\mathbf{Q}' ((\mathbf{K}')^\top \mathbf{V})), \quad \widehat{\mathbf{D}} = \text{diag}(\mathbf{Q}' ((\mathbf{K}')^\top \mathbf{1}_L)).$$

It uses low-rank approximation to query and key matrices: $\mathbf{Q}' \in \mathbb{R}^{T \times r}$, $\mathbf{K}' \in \mathbb{R}^{r \times T}$



Choromanski, Krzysztof Marcin, et al. *Rethinking Attention with Performers*. ICLR. 2020.

Exponential kernel decomposition

Kernel representation

$$\exp(q_i^T k_j) = \phi(q_i)^T \phi(k_j) = \mathbb{E}(\phi(q_i)^T \phi(k_j))$$

Feature vector

$$\phi(\mathbf{x}) = \frac{h(\mathbf{x})}{\sqrt{m}} (f_1(\omega_1^\top \mathbf{x}), \dots, f_1(\omega_m^\top \mathbf{x}), \dots, f_l(\omega_1^\top \mathbf{x}), \dots, f_l(\omega_m^\top \mathbf{x}))$$

$$\mathbf{x}, \mathbf{y} \in \mathbb{R}^d, \mathbf{z} = \mathbf{x} + \mathbf{y}$$

Our kernel

$$\Lambda = \exp\left(-\frac{\|\mathbf{x}\|^2 + \|\mathbf{y}\|^2}{2}\right)$$


$$\text{SM}(\mathbf{x}, \mathbf{y}) = \exp\left(\frac{\|\mathbf{x}\|^2}{2}\right) K_{\text{gauss}}(\mathbf{x}, \mathbf{y}) \exp\left(\frac{\|\mathbf{y}\|^2}{2}\right)$$

$$\text{SM}(\mathbf{x}, \mathbf{y}) = \mathbb{E}_{\omega \sim \mathcal{N}(0, \mathbf{I}_d)} \left[\exp\left(\omega^\top \mathbf{x} - \frac{\|\mathbf{x}\|^2}{2}\right) \exp\left(\omega^\top \mathbf{y} - \frac{\|\mathbf{y}\|^2}{2}\right) \right] = \Lambda \mathbb{E}_{\omega \sim \mathcal{N}(0, \mathbf{I}_d)} \cosh(\omega^\top \mathbf{z})$$


Pathfinder problem

For an image reshaped to the image find out, if two points are connected by a dashed line or not.

- Pathfinder: image 32x32 -> sequence 1024
- Pathfinder-X: images 128x128 -> sequence 16K, extreme length (c.t. 1K limit for GPT-3)



Positive example



Negative example

All transformers fail for Path-X problem

| Model | ListOps | Text | Retrieval | Image | Pathfinder | Path-X | Avg |
|-----------------|--------------|--------------|--------------|--------------|--------------|--------|--------------|
| Transformer | 36.37 | 64.27 | 57.46 | 42.44 | 71.40 | FAIL | <u>54.39</u> |
| Local Attention | 15.82 | 52.98 | 53.39 | 41.46 | 66.63 | FAIL | 46.06 |
| Sparse Trans. | 17.07 | 63.58 | 59.59 | 44.24 | 71.71 | FAIL | 51.24 |
| Longformer | 35.63 | 62.85 | 56.89 | 42.22 | 69.71 | FAIL | 53.46 |
| Linformer | 35.70 | 53.94 | 52.27 | 38.56 | <u>76.34</u> | FAIL | 51.36 |
| Reformer | 37.27 | 56.10 | 53.40 | 38.07 | 68.50 | FAIL | 50.67 |
| Sinkhorn Trans. | 33.67 | 61.20 | 53.83 | 41.23 | 67.45 | FAIL | 51.39 |
| Synthesizer | <u>36.99</u> | 61.68 | 54.67 | 41.61 | 69.45 | FAIL | 52.88 |
| BigBird | 36.05 | 64.02 | <u>59.29</u> | 40.83 | 74.87 | FAIL | 55.01 |
| Linear Trans. | 16.13 | 65.90 | 53.09 | 42.34 | 75.30 | FAIL | 50.55 |
| Performer | 18.01 | <u>65.40</u> | 53.82 | <u>42.77</u> | 77.05 | FAIL | 51.41 |
| Task Avg (Std) | 29 (9.7) | 61 (4.6) | 55 (2.6) | 41 (1.8) | 72 (3.7) | FAIL | 52 (2.4) |

Experimental results for Long-Range Arena Benchmark


Tay, Yi, et al. "Long range arena: A benchmark for efficient transformers." *arXiv preprint arXiv:2011.04006* (2020).

Efficiently Modeling Long Sequences with Structured State Spaces


Gu, Albert, Karan Goel, and Christopher Re. Efficiently Modeling Long Sequences
with Structured State Spaces. *ICLR*. 2021.

Paradigms for sequence modelling


Continuous-time Model




Recurrent Neural Net.



Convolutional Neural Net.



Transformer



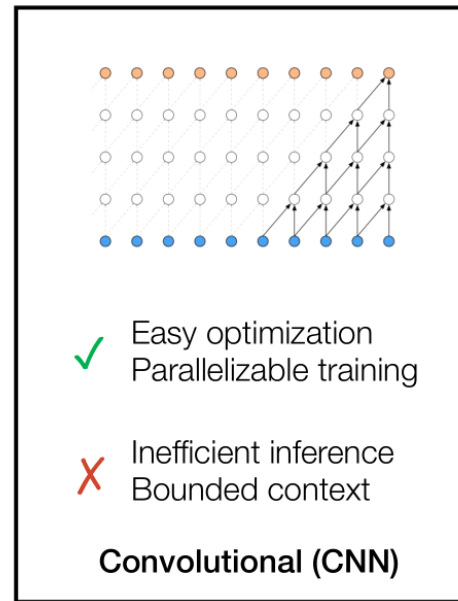
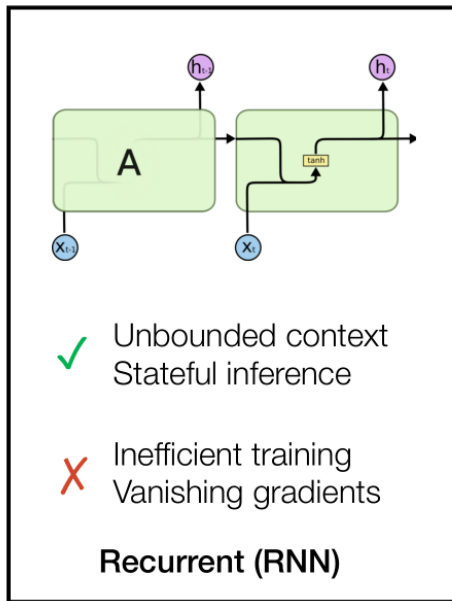

Deep Sequence Model
Sequence-to-sequence map

Sequence Model Layer

(batch, length, dim)

(batch, length, dim)

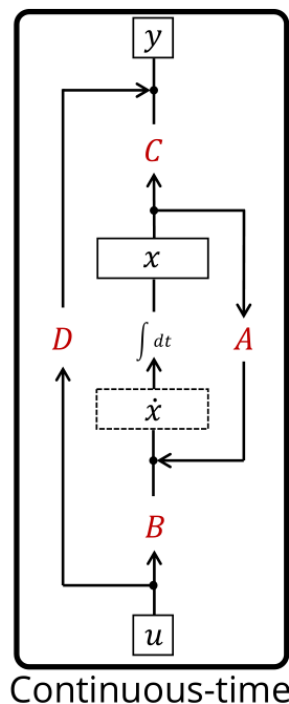
Pros and cons for each paradigm



Existing model families have clear tradeoffs
All struggle with long-range dependencies (LRD)

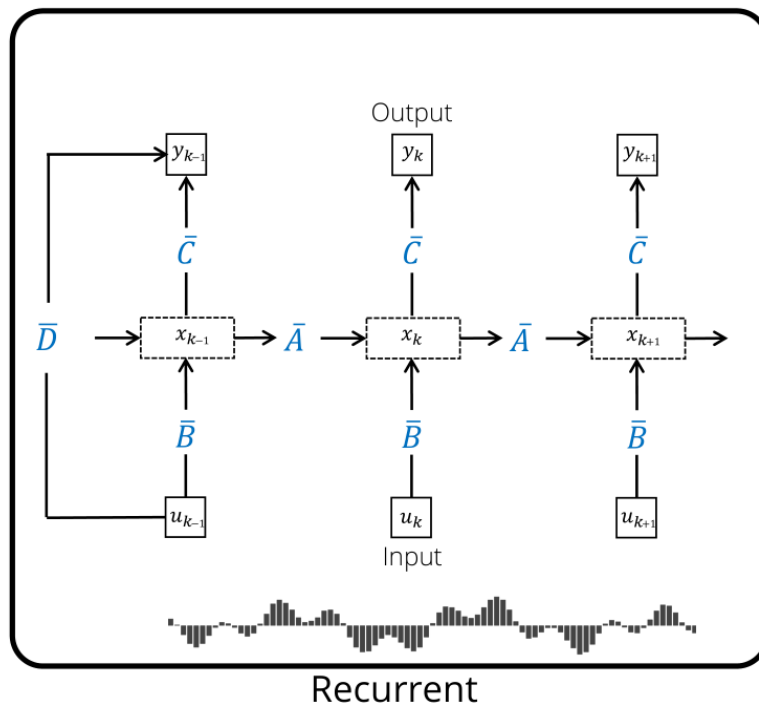
Gu, Albert, Karan Goel, and Christopher Re. Efficiently Modeling Long Sequences with Structured State Spaces. *ICLR*. 2021.

Structured state-space models




Discretize

Δt



or



They are all similar! And you will soon know how...

Long-range arena benchmark


Benchmark spanning text, images, symbolic reasoning (length 1K-16K)

| Model | LISTOPS | TEXT | RETRIEVAL | IMAGE | PATHFINDER | PATH-X | AVG |
|-----------------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|
| Random | 10.00 | 50.00 | 50.00 | 10.00 | 50.00 | 50.00 | 36.67 |
| Transformer | 36.37 | 64.27 | 57.46 | 42.44 | 71.40 | X | 53.66 |
| Local Attention | 15.82 | 52.98 | 53.39 | 41.46 | 66.63 | X | 46.71 |
| Sparse Trans. | 17.07 | 63.58 | 59.59 | 44.24 | 71.71 | X | 51.03 |
| Longformer | 35.63 | 62.85 | 56.89 | 42.22 | 69.71 | X | 52.88 |
| Linformer | 35.70 | 53.94 | 52.27 | 38.56 | 76.34 | X | 51.14 |
| Reformer | <u>37.27</u> | 56.10 | 53.40 | 38.07 | 68.50 | X | 50.56 |
| Sinkhorn Trans. | 33.67 | 61.20 | 53.83 | 41.23 | 67.45 | X | 51.23 |
| Synthesizer | 36.99 | 61.68 | 54.67 | 41.61 | 69.45 | X | 52.40 |
| BigBird | 36.05 | 64.02 | 59.29 | 40.83 | 74.87 | X | 54.17 |
| Linear Trans. | 16.13 | <u>65.90</u> | 53.09 | 42.34 | 75.30 | X | 50.46 |
| Performer | 18.01 | 65.40 | 53.82 | 42.77 | 77.05 | X | 51.18 |
| FNet | 35.33 | 65.11 | 59.61 | 38.67 | <u>77.80</u> | X | 54.42 |
| Nyströmformer | 37.15 | 65.52 | <u>79.56</u> | 41.58 | 70.94 | X | 57.46 |
| Luna-256 | 37.25 | 64.57 | 79.29 | <u>47.38</u> | 77.72 | X | <u>59.37</u> |
| S4 | 58.35 | 76.02 | 87.09 | 87.26 | 86.05 | 88.10 | 80.48 |

Path-X



(a) A positive example.



(b) A negative example.

State Space Models (SSM)

Input → State

$$x'(t) = Ax(t) + Bu(t)$$

$$y(t) = Cx(t) + Du(t)$$

Parameters

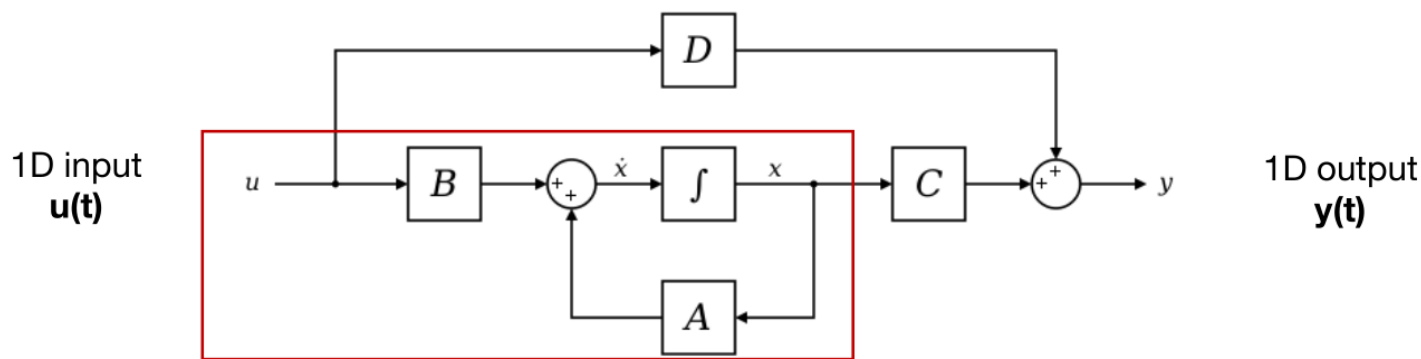
$$A \in \mathbb{R}^{N \times N}$$

$$B \in \mathbb{R}^{N \times 1}$$

$$C \in \mathbb{R}^{1 \times N}$$

$$D \in \mathbb{R}^{1 \times 1}$$

Function-to-function map $u(t) \mapsto y(t)$



State Space Models (SSM)

Parameters

$$\mathbf{A} \in \mathbb{R}^{N \times N}$$

$$\mathbf{B} \in \mathbb{R}^{N \times 1}$$

$$\mathbf{C} \in \mathbb{R}^{1 \times N}$$

$$\mathbf{D} \in \mathbb{R}^{1 \times 1}$$

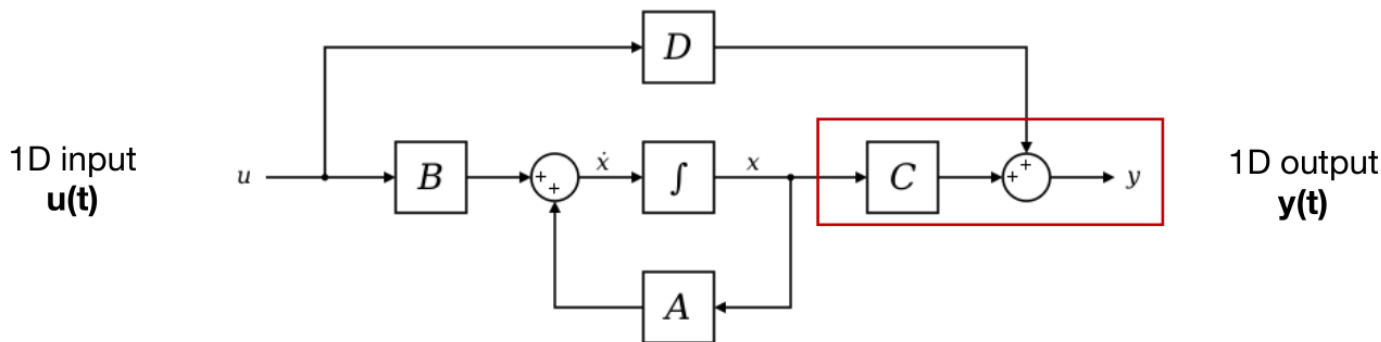
State \rightarrow Output

$$x'(t) = \mathbf{A}x(t) + \mathbf{B}u(t)$$

$$y(t) = \mathbf{C}x(t) + \mathbf{D}u(t)$$

Function-to-function map

$$u(t) \mapsto y(t)$$



Computing with SSMs: Recurrent View

$$x'(t) = \mathbf{A}x(t) + \mathbf{B}u(t)$$

$$y(t) = \mathbf{C}x(t) + \mathbf{D}u(t)$$

Parameters

$$\mathbf{A} \in \mathbb{R}^{N \times N}$$

$$\mathbf{B} \in \mathbb{R}^{N \times 1}$$

$$\mathbf{C} \in \mathbb{R}^{1 \times N}$$

$$\mathbf{D} \in \mathbb{R}^{1 \times 1}$$

$$\Delta \in \mathbb{R}$$

1. Discretize

$$\overline{\mathbf{A}} = \mathbf{I} + \Delta \mathbf{A}$$

2. Recurrent "hidden state"

$$x_k = \overline{\mathbf{A}}x_{k-1} + \overline{\mathbf{B}}u_k$$

3. Out projection

$$y_k = \overline{\mathbf{C}}x_k + \overline{\mathbf{D}}u_k$$

Can be computed with linear recurrence, similar to RNNs

Computing with SSMs: Convolution View

$$x_k = \overline{A}x_{k-1} + \overline{B}u_k$$

$$y_k = \overline{C}x_k$$

Can explicitly unroll the linear recurrence in closed form

$$\begin{aligned} x_0 &= \overline{B}u_0 & x_1 &= \overline{A}\overline{B}u_0 + \overline{B}u_1 & x_2 &= \overline{A}^2\overline{B}u_0 + \overline{A}\overline{B}u_1 + \overline{B}u_2 & \dots \\ y_0 &= \overline{C}\overline{B}u_0 & y_1 &= \overline{C}\overline{A}\overline{B}u_0 + \overline{C}\overline{B}u_1 & y_2 &= \overline{C}\overline{A}^2\overline{B}u_0 + \overline{C}\overline{A}\overline{B}u_1 + \overline{C}\overline{B}u_2 & \dots \end{aligned}$$

Computing with SSMs: Convolution View

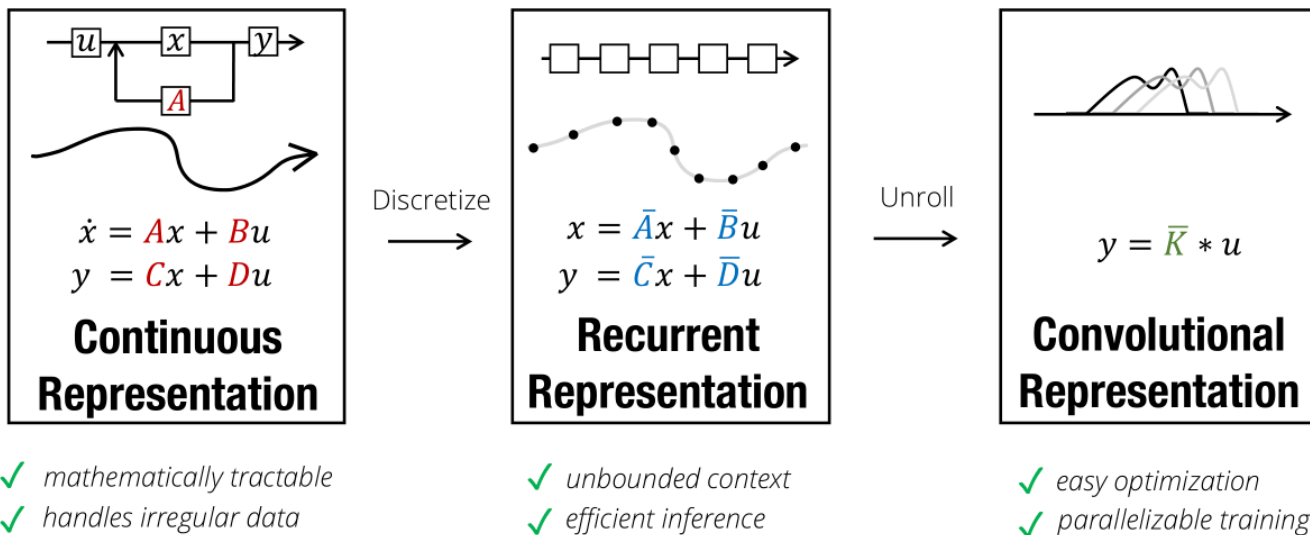
$$y_k = \overline{CA}^k \overline{B} u_0 + \overline{CA}^{k-1} \overline{B} u_1 + \cdots + \overline{CAB} u_{k-1} + \overline{CB} u_k$$

$$\overline{K} \in \mathbb{R}^L := (\overline{CB}, \overline{CAB}, \dots, \overline{CA}^{L-1} \overline{B})$$

$$y = \overline{K} * u$$

Can be computed with **convolutions**, similar to CNNs

Summary: Properties of SSMs



Conclusions

- We can improve Transformers memory and computational requirements with some heuristics and different types of attention
- The main ideas are:
 - combine different sparse attention mechanisms
 - bucket processing
- But there are some limitations in terms of reducing number of parameters and memory as we have large window sizes and large number of global attention nodes